

# Introduction

This page describes the general RA guidelines we employ for our projects. The purpose of the guidelines is to make it easier to focus on doing replicable and properly documented work. The guidelines draw heavily on

- *Code and Data for the Social Sciences: A Practitioners Guide*. Gentzkow and Shapiro (2014). [Link](#)
  - The document contains a general introduction to reproducibility challenges that many social science researchers have faced and how the two authors have tried to solve them in their lab. Emphasizes coding conventions, directories, and version control.
- *Lab Manual*. Gentzkow and Shapiro. [Link](#)
  - This supplements their (2014) document. It includes information on workflows, coding, data handling, and paper and slide production. We generally follow their workflow approaches (e.g., using GitHub issues to assign and resolve tasks).

How to get started:

- Read [Gentzkow and Shapiro \(2014\)](#).
- Read this document.
  - You do not have to follow links the first time you read the document. They are there for reference.
- Install the necessary software on your local computer as described under `Software/Required`.
- Set up git on your computer and a GitHub account as described under `Version Control`.
- Ask your supervisor to add you to the relevant projects on GitHub and Zotero.

After you have finished these steps, we will assign tasks to you via the Github project(s).

## Workflow

### Tasks

- We follow Gentzkow and Shapiro's workflow [approach](#).
- Tasks will be specified on GitHub under the relevant project as `issues`. We do this to keep track of open tasks, task notes and questions, and outputs. You can familiarize yourself with GitHub issues [here](#).
  - A supervisor will add you to the relevant project(s) on GitHub.
- Each task will contain
  - A description.
  - A set of outcomes.
  - A task supervisor.
  - One or more task assignees.
- When working on a task
  - Keep documentation of your work. We suggest having a `running_notes_(your initials).md` document where you store your thoughts. A good practice is to add headlines with dates so that it becomes easier to go back to find thoughts related to a task you previously worked on.

- Asking questions
  - We encourage you to work independently but ask questions when you realize you are stuck or something seems unclear. We all get stuck. And it happens particularly often when we start working with the administrative data, so do come around and ask. Our experience is that RAs who ask questions early on are more productive in the long run.
  - If you ask about clarifications or questions related to a task in person, add a note about the questions and answers to the GitHub issue so we can track the progress.
- A task is closed by the task supervisor when
  - The relevant outcomes have been created/reached.
  - The task assignee has written a reply to the GitHub issue on how they completed the task and where relevant outcome files are located (e.g., the code file cleaning a bit of data or the note summarizing results).
  - The task supervisor agrees that the task is completed.

## Reporting and notes

- We write notes and documents in markdown format with files ending in `.md` unless another format is required.
  - Markdown documents can easily be compiled into Word, PDF (via LaTeX), HTML, beamer PDF slides, or other formats using `pandoc` or `quarto`.
  - Searching (and replacing) across multiple markdown files for content using typical text editors is easy.
  - Markdown files can be edited using most text editors. We suggest VSCode and Obsidian.
  - [Introduction to basic markdown syntax](#) written by the developers behind the original markdown language (it comes in many flavors).
- We keep personal running notes documents.
  - These typically contain thoughts and drafts for notes and tasks.
  - The document will typically be named `running_notes_` and end in our initials, e.g., `running_notes_je.md`.

## References

- We use Zotero to maintain shared libraries with project references.
  - A supervisor will add you to the relevant project(s) on Zotero.
- We use the `betterbiblatex` extension for Zotero to export `.bib` files to projects or to Overleaf.
- When adding a new reference to a Zotero project collection, pin the bibtex key.
  - Right-click the reference and select `better bibtex -> pin citekey`.
  - Remember to set up the `better bibtex` extension to use the correct citekey structure. See [Software / Required / Zotero](#).

## Writing papers

- We use Overleaf to write papers and paper drafts in LaTeX format unless otherwise specified. Overleaf allows us to
  - Work on the paper at the same time
  - Integrate references from shared Zotero libraries

# Version Control

- We use the [Git version control system](#) to track changes to our notes and code. Git allows you to add, delete, or modify files, mark them as changed (committing), and finally add them to the project (pushing to the remote repository). All changes to files will then appear in the project Git history.
- We use GitHub to host our remote repositories for *non-sensitive* project files.
  - You can keep the local non-sensitive project files where you prefer.
  - You will pull and push changes to this remote directory to update general project files.
  - A supervisor will add you to the relevant project(s) on GitHub.
- We use a local Git repository for content stored on secure servers where we cannot access Github.
- A basic workflow for using Git is to Pull the latest version of project files from the remote repository.
  1. Pull latest updates from the remote repository.
  2. Make changes to the files you are working on.
  3. Commit your changes.
  4. Add a reference to the task you are working on in the commit message if relevant. You do this by adding adding a reference to the task-number, e.g., `Task24 -` , in front of the commit message.
  5. Pull the latest updates from the remote repository and resolve any conflicts.
  6. Push your changes to the remote repository.
- Guides
  - Basic (takes about 20 minutes total and well worth it): [Git Handbook](#), [Understanding the GitHub Flow](#), [Mastering Issues](#), [Mastering Markdown](#)
  - Detailed: [Pro Git](#), [Version Control with Git](#), Chapters 4-9
- Setup: (based on <https://github.com/gslab-econ/lab-manual/wiki/Setup>)
  - [Create a GitHub account and install the Git desktop/command line clients](#).

## Project structure

### Project storage

- We generally store non-sensitive project files in private GitHub repositories (one per project). Using Github repositories ensures all project participants can access the relevant files.
  - It also makes it easy to track file changes and revert to previous versions if necessary.
  - We use the `git` version control system to pull and push files to the Github repositories.
  - Supervisors will add you to the relevant project(s) on GitHub.
- We store sensitive files on a secure KU server.
  - Supervisors can give you access to relevant folders.
- GDPR-sensitive microdata is stored on a secure server hosted at Statistics Denmark. This includes administrative data from Statistics Denmark and the Ministry of Education we use in our projects.
  - Documentation on Statistics Denmark's researcher data access and storage is available [here](#).
  - If you work with sensitive microdata, your project supervisor will help you set up access.
  - Before getting access, you are required to read and sign the internal UCPH ECON and Statistics Denmark guidelines on working with sensitive data.

- We follow the internal UCPH ECON guidelines on working with sensitive data, including what information can be downloaded from the secure server.

## Directories

- Our projects are typically distributed locally and on a secure server. The local project might contain notes, literature reviews, code that doesn't need to run on a secure server, and paper drafts. The secure server project location will exist when the project requires restricted access data (e.g., from Statistics Denmark).
- We generally apply the rules from Gentzkow and Shapiro (2014), chapter 4:
  1. Separate directories (folders) by function.
  2. Separate files into inputs and outputs (and temporary files)
  3. Make directories (folders) portable.
- Each (sub)component of a project should have its separate folder. A project with a literature review, presentation files, and a paper (draft) should contain at least those folders.

```
lit_review/
presentations/
paper/
```

- All components containing code should have at least a `code`, `temp`, and `out` folder. For example, assume that the simple project contains a simulation exercise written in R showing the consistency of an econometric estimator. The code file outputs the graph

`simulate_estimator_consistency_distribution.PDF`. The folders could look like

```
consistency_simulation/
  code/
    simulate_estimator_consistency.R
  temp/
  out/
    simulate_estimator_consistency_distribution.PDF
lit_review/
presentations/
paper/
```

- We store raw data in a separate folder. Suppose we have more than one raw data set, for example, from Statistics Denmark and the Ministry of Education. We then separate them into subfolders with meaningful names and possibly a date of compilation so we can keep track of versions.

```
buildraw/
  dst/
    dat1.sas7bdat
    dat2.sas7bdat
consistency_simulation/
  code/
    simulate_estimator_consistency.R
  temp/
  out/
    simulate_estimator_consistency_distribution.PDF
lit_review/
```

```
presentations/  
paper/
```

## Data formats

- When possible, store data in `.parquet` format.
  - We often work with large administrative data files that can take up many GB of space. We prefer the parquet format to reduce our server footprint and increase IO speed.
- We can import and export `.parquet` files using `R` and `python`.
  - [Guide to using parquet with R](#) by Apache.
  - [Guide to using parquet with R](#) by Hadley Wiggins.
- In R, we will often use the `rio` package to import/export a dataset:

```
library(rio)  
dat = import("builddata/out/clean_bef.parquet")  
dat ► export("builddata/out/clean_bef.parquet")
```

## R projects

- We use R projects to help RStudio determine where to run our R scripts.
  - The R-project file must be located at the root of the project directory.
  - Read about working with scripts and R projects [here](#).

## Code Conventions

### General

- Name scripts by what they do.
  - *Example:* Assume we have written some code that cleans the raw `BEF` register data for use in subsequent analyses. The file (sh)could be named `clean_bef.R`.
- Script outputs must contain the name of the producing script and be informative about the content.
  - *Example:* Assume the file `descriptives_main_sample.R` outputs two summary tables in LaTeX format. One is a balance table with means and differences in means between treatment and control groups, and one contains general summary statistics for the full sample. These (sh)could be named `descriptives_main_sample_balance.tex` and `descriptives_main_sample_summary.tex`.
- No line of code should be more than 100 characters long. All languages we work in allow you to break a logical line across multiple lines on the page (e.g., using `///` in Stata or `...` in Matlab). You may want your editor to show a “margin” of 100 characters.
- Functions should not typically be longer than 200 lines.

## R

- We follow [Google's R Style Guide](#)
- Exceptions to style guide:
  - Do not use dots, separate using underscores, and keep lowercase. Example: `.CalcMeans()` should be `calc_mean()`.

- Use the `rio` [package](#) for data import/export.
  - It supports many file types and generally uses the most efficient IO tool for importing/exporting the file format.
- Use the `data.table` [package](#) for data wrangling when possible.
  - [Introduction to data.table](#) from the authors of the package.
  - [Introduction for Stata users](#)
  - [data.table chapter](#) in Introduction to Data Science.
- Use the `fixest` [package](#) by Laurent Berge for estimating most types of statistical models, particularly linear and IV models with fixed effects, when possible.
  - It provides estimation tools typically much faster than other options in R and Stata.
  - Linear, fixed effects, and 2SLS models can be estimated via `feols()`.
  - [Documentation](#).
- We typically use the `modelsummary` [package](#) for summarizing regression results and creating summary statistics tables outputted to latex or markdown format.
  - An introduction is available [here](#).
  - When adding footnotes to `modelsummary()` tables, use the `footnote()` function from the `kableExtra` package with the options `escape = F` and `threeparttable = T`.
- We use the [here](#) [package](#) when specifying paths to ensure all paths are read relative to the project folder.

## Guides to get started with R

- How to get started:
  - Read parts 1 and 2 of Hans Henrik Sievertsen's [Introduction to R](#) and solve the associated exercises.
    - The introduction focuses on the basics of R and the `tidyverse` packages.
  - Read through Atrebas' [introduction to using the data.table package](#).
    - The document describes basics like viewing data, subsetting, creating new variables, and using the `.SD` capability.
  - Go through the `r_introduction.qmd` quarto document that you can find on this Github page.
- Other great guides
  - [R for Data Science \(2e\)](#) is an online and free goto reference for getting started with simple and more advanced R, including data IO with [arrow](#), writing [functions](#), and using [quarto](#) to communicate results.
    - Each chapter comes with great exercises.
    - The main author, Hadley Wickham, has been a driving force in developing R packages, including a majority of the `tidyverse` package since the 2000s.
  - Hans Henrik Sievertsen's [Applied Econometrics with R](#) introduces using R for applied econometrics, including data cleaning, visualization, descriptive statistics, and regression analysis (using `feols` and `modelsummary`). This guide is great for when you've gotten the hang of basic R.
  - Hans Henrik Sievertsen's [Interactive introduction to R](#). It introduces basic data handling (loading data, modifying and merging datasets) and plotting.
- Remember, ChatGPT often gives great solutions to coding problems!

## Automation

- We automate everything that can be automated. This implies writing scripts to do all data cleaning, analysis, and table formatting, and using build tools to run these scripts in the correct order.
- We use the build tool `make` to run all project code.
  - Generally, we want to be able to delete all files in output folders, type `make all` on the command line, and have all output files reappear just as they were before.

## Build tool - `make`

- Build tools generally help run your project code in the correct order based on a set of rules that specify output targets and inputs, including code files. Some examples of more advanced build tools include `snakemake`, and `cmake`. I have generally found that these modern build tools are unavailable on the secure servers hosted at Statistics Denmark, where a substantial part of my project code resides.
- We, therefore, use the slightly more archaic `make`.
- Documentation for `make` is available [here](#).
- Helpful introductions to `make` for data analysis
  - [Automation and Make](#) by Software Carpentry.
  - [Makefiles for R/Latex projects](#) by Rob Hyndman.
  - [Minimal make](#) by Karl Broman. Runs a couple of R scripts and creates a latex compiled PDF paper with the resulting figures.
  - [GNU Make for Reproducible Data Analysis](#) by Zachary Jones.
- An important feature of `make` is that it compiles or runs project code based on a general recipe, the `makefile`. The `makefile` consists of *targets*, *dependencies*, and *commands*, which together defines *rules*. A `makefile` can contain multiple rules. These rules can be linked, for example, if a rule uses the target of another rule as a dependency.
  - *Target*: The output file or goal you want to achieve. E.g.,  
`builddata/out/clean_data.parquet`
  - *Dependencies*: Files or targets that must be up-to-date before executing the target's commands. This will typically include the code file you want to run, and the data it uses. E.g.,  
`builddata/code/clean_data.R`, and `buildraw/out/rawdata.csv`.
  - *Command*: The command(s) that `make` will execute to create or update the target. `make` basically runs from the command line meaning that you must either use a command that is available from the CLI or specify the full path to the program. If, for example, you want to run an R script, you can use the CLI command `rscrip` when this is installed.

```
#target: dependencies
#      command

builddata/out/clean_data.parquet: builddata/code/clean_data.R buildraw/out/rawdata.csv
    rscrip builddata/code/clean_data.R
```

- Make `make` run your command
  1. Navigate a command line tool, such as `cmd` to the project folder where the `makefile` is located.
  2. In the command line, type `make target`, where `target` is the file you want to build. E.g., `make builddata/out/clean_data.parquet`.
- How `make` runs your command:
  - When you tell `make` to create a target, `make` first checks the dependencies for that rule. If a dependency is itself a target from another rule, `make` moves back to this previous rule. This

continuous until `make` finds the antecedent dependencies.

- `make` then checks the timestamps of antecedent dependencies. If the dependencies have not changed since the target was last created, `make` won't re-run the commands for that target. If the target does not exist, the command is always run.
- `make` will move through the linked set of dependencies, running the commands from rules where dependencies have changed since the target was last created.
- `make` will let you know before it tries to run if a dependency does not exist and there exists no rule to create it.
- Installation:
  - `make` is usually pre-installed on Unix-based systems (like macOS and Linux). For Windows, you can install it through tools like Cygwin, GNUWin32, or `rtools`.
- Important syntax notes:
  - **Indentation**: Make sure to use a tab, not spaces, for indentation in the Makefile.
  - **Multiple Languages**: If your workflow involves Stata or SAS scripts, you can include them in the Makefile just like R scripts. For instance, `stata -b do my_analysis.do` for a Stata script.
- Line splitting with many dependencies
  - If your file contains many dependencies it can be useful to split the dependency list over multiple lines. You can do this by writing `\` and continuing the content on the next line after an indentation. If you include anything, include a space, after the backslash, `make` will throw an error.

```
builddata/out/clean_data.parquet: \  
    builddata/code/clean_data.R \  
    builddata/out/rawdata.csv  
    rscript builddata/code/clean_data.R
```

- Many targets
  - A file may create multiple outputs, such as regression tables. To specify this, you simply add all targets to the left of `:`

```
builddata/out/clean_data1.parquet builddata/out/clean_data2.parquet : \  
    builddata/code/clean_data.R \  
    builddata/out/rawdata.csv  
    rscript builddata/code/clean_data.R
```

- Automatic variables
  - `make` allows you to use [automatic variables](#) in rules.
    - `$@`: The name of the target
    - `$<`: The name of the first prerequisite
    - `$$`: The names of all the prerequisites, with spaces between them

```
builddata/out/clean_data.parquet: \  
    builddata/code/clean_data.R \  
    builddata/out/rawdata.csv  
    rscript $<
```

- Creating variables
  - Creating variables can be useful, for example, for creating build rules or listing all targets you want to build.



- Variables can be assigned with `:=` . E.g., `var := name1`
- Variables can be called using `$(variablename)`
- You can append to a variable by adding `$(variablename)` on the right hand side of the assignment name; `var := $(var) name2` .

```
R := C:\Users\bxn825\scoop\shims\rscript.exe
```

```
builddata/out/clean_data.parquet: \
    builddata/code/clean_data.R \
    buildraw/out/rawdata.csv
$(R) $<
```

- Creating build rules with many targets (phony targets)
  - The rule `all: $(targets)` is often used to run all rules you want in a project, when the variable `$(targets)` contains a set of targets in your project.
  - You can use `.PHONY` to explicitly declare `all` (and other non-file targets) as a phony target. This tells `make` that this target isn't a file but rather a label for a recipe to be executed.

```
R := C:\Users\bxn825\scoop\shims\rscript.exe
```

```
$(targets) := builddata/out/clean_data.parquet
builddata/out/clean_data.parquet: \
    builddata/code/clean_data.R \
    buildraw/out/rawdata.csv
$(R) $<
```

```
$(targets) := $(targets) builddata/out/clean_data.parquet
builddata/out/clean_data.parquet: \
    builddata/code/clean_data.R \
    buildraw/out/rawdata.csv
$(R) $<
```

```
.PHONY: all
```

```
all: $(targets)
```

## Software

### Required

If you work on a windows machine, do consider installing the software using `scoop` . We have created a PowerShell executable file that can install all required software in one go for you. To run it:

1. first [install scoop](#)
2. Download the file `scoop_install_required_software.ps1` from the `_setup` folder in this repository.
3. Right-click it and choose "run with powershell".

The file then installs all the software. Remember you still have to install extension within each piece of software manually.

## List of required software

- Git
  - [Git](#) is a command-line version control software. See more under [Version Control](#)
  - Installation
    - [Create a GitHub account and install the Git desktop/command line clients.](#)
    - Installation via Scoop on Windows machines: `scoop install git`
- VSCode
  - [VSCode](#) is a general text editor you can use to edit markdown, Python, R, LaTeX, and many other file types.
  - It has excellent Git version control features included
  - Installation options:
    - Follow Jeppe Druedahl's [guide to install VSCode](#).
    - Install using Scoop on a Windows machine with `scoop bucket add extras; scoop install vscode` (see [Software/Suggestions](#) on using ).
  - You can add functionality by installing extensions. Some extensions you'll likely want to install:
    - Git History
    - GitHub Copilot (free for academic users)
    - LaTeX workshop
    - Markdown All in One
    - Python
    - Quarto
    - R (consider using this together with `radian`)
    - Stata Enhanced
    - vscode-pandoc
  - Guides:
    - [Markdown and Visual Studio Code](#)
  - We use `pandoc` or `quarto` to convert markdown files into Word, PDF (via LaTeX), beamer slides, or HTML files.
- R, RStudio, rtools
  - [R](#) is an open-source statistical programming language that has gotten a lot of traction in the Econometrics community. Most new developments in Econometrics are likely to arrive to R at the same time or prior to, e.g., Stata.
  - [Rstudio](#) is an editor specialized for R coding. We typically use RStudio whenever we edit R code.
  - [rtools](#) is a set of software tools R will need to compile some packages, including the `arrow` package we use for parquet format IO.
  - Installation options:
    - [Installation guide for R and RStudio] by Posit (the developers behind RStudio).
    - [Installation guide for rtools](#) via CRAN.
    - Installation via Scoop on Windows machines: `Scoop bucket add r-bucket https://github.com/cderv/r-bucket.git; scoop install r; scoop install rstudio; scoop install rtools`
    - Guides to get started using R
      - [4h R crash course](#) by Hans H. Sievertsen. With Economics related examples
      - [R for Data Science](#) by Hadley Wickham and Garrett Grolemund.

- Tinytex
  - [Tinytex](#) is a lightweight LaTeX distribution that allows you to compile LaTeX documents to PDF.
  - Installation options:
    - [Guide](#) from the developers.
    - Installation via Scoop on Windows machines: `scoop install tinytex`
- Pandoc
  - [Pandoc](#) is an open-source command line tool that can be used to convert between many different text formats. We will typically use it to convert between markdown and PDF/word documents.
  - Installation options:
    - [Guide](#) from the developers. Installation via Scoop on Windows machines: `scoop install pandoc`
- Quarto
  - [Quarto](#) is a piece of software developed by Posit that allows you to write quarto documents containing both markdown text and integrated R, Python, or some other relevant code. The documents can be edited in RStudio or VSCode. Quarto documents are beneficial when some code requires extensive documentation.
  - Installation options:
    - [Guide](#) from the developers.
    - Installation via Scoop on Windows machines: `scoop install quarto`
- Zotero
  - [Zotero](#) is an open-source reference manager.
  - Allows easy export of bibtex reference libraries that can be cited in markdown and latex documents.
  - Installation options:
    - [Guide](#) from the developers.
    - Installation via Scoop on Windows machines: `scoop install zotero`
  - Install connectors:
    - [Installation guide](#) from the Zotero developers.
  - Install relevant extensions
    - [better bibtex](#)
      - Install using the guide at the link.
      - Set bbt's citation key formula to `authEtAl + year + shorttitle(3,3)` by going to preferences → Better BibTeX → Open Better BibTeX Preferences → Citation Key, and copy the formula in.

## Suggestions

- Scoop
  - A Powershell tool for Windows that helps you install and update software updated.
  - Installation options: - [Guide](#) from the developers.
  - Working on university-provided IT equipment can give update and installation problems if you do not have administrator rights over the computer. This can be circumvented by ensuring that all (or most) programs are installed in your own user path. Scoop does this.
  - We presently use the Powershell tool `scoop` to manage the installation of most software on my system. `scoop` uses recipes created by others to install (often) the latest versions of

programs.

- Installing a program with `scoop` is as simple as `scoop install program`.
- Updating a program with `scoop` is as simple as `Scoop update program`. Type `scoop update *` to update all installed programs.
- Scoop searches for install recipes in buckets. Buckets can be added by typing `scoop bucket add ...` in Powershell. Examples of useful buckets include `extras`, `nerd-fonts`, and `r-bucket`. You can find an example of installing and using a `scoop` to set up a new Windows machine at [https://github.com/EriksenJ/\\_setup](https://github.com/EriksenJ/_setup).
- Obsidian
  - A markdown-based digital note editor with latex compilation to PDF ready out of the box.
  - Some suggested plugins
    - LaTeX Suite
    - Linter
    - Obsidian Git
    - Templater
    - Zotero Integration
  - Install with `scoop` on windows: `scoop bucket add extras; scoop install obsidian`
- Linh T. Tô has a great set of (free) resources on her [website](#).